

---

# LOCK FREE MULTIDIMENSIONAL RANGE SEARCH

---

**Aman Aggarwal**  
2018327  
IIIT  
Delhi  
aman18327@iiitd.ac.in

**Prasham Narayan**  
2018359  
IIIT  
Delhi  
prasham18359@iiitd.ac.in

## ABSTRACT

Multi-Dimensional Range Search is a fundamental problem in many application domains dealing with multidimensional data. R-tree is an ideal candidate for the multidimensional range search problem. However, there is no existing implementation of R-tree supporting concurrent operations. This paper introduces the lock-free concurrent R-tree, which implements an abstract data type (ADT) that provides the operations Add, Remove, RangeSearch. The operations in the Lock Free R-tree use single-word read and compare-and-swap (CAS) atomic primitives, which are readily supported on available multi-core processors. A lock-based implementation is also provided as a benchmark.

**Keywords** Lock-Free · Lock-Based · R-tree · Range Search

## 1 Introduction

We propose using an external 2-way R-tree to implement the ADT = {Add, Delete, RangeSearch} in a 2-dimensional setting. An R-Tree is a tree-based data structure that can store spatial data indexes. Each entry in an internal node of an R-tree corresponds to MBR (Minimum bounding rectangle), which represents the smallest rectangle containing all its child nodes. R-trees use this concept of MBR to group nearby objects, which helps in spatial search. Each entry in the leaf nodes corresponds to a spatial point.

The main structure of our Lock free R tree is adapted from its sequential implementation. In the sequential implementation, we defined all the classes with its internal structure. Using the sequential implementation we also created a lock-based implementation of R tree. This lock-based implementation provided further insights on additional cases that are needed to be handled in concurrent setting.

Moreover, we are also planning to compare the performance of lock-based implementation of R tree with its lock free implementation. Github Link - [github.com/aman-agg/CLDS-Project](https://github.com/aman-agg/CLDS-Project)

## 2 Abstract Data Type

Multi Dimensional Range Search Tree (MDRST) = {Add, Delete, RangeSearch}

- **Add** - adds point  $p$  to the tree
- **Delete** - deletes point  $p$  from the tree
- **RangeSearch** - Given a 2-dimensional range, return points lying inside that range.

### 2.1 Invariants

- Two way R tree i.e. it has two children (leftChild and rightChild).
- There will be atleast a single entry in a node.
- All internal nodes will have both the entries.
- MBR criteria i.e. parent node's MBR will envelop the MBR of its children.
- Tree can be imbalanced.

### 3 Structure of Components

We will be using a 2-way R-tree. In a 2-way R-tree, each node consists of 2 entries (left entry and right entry). Each node contains a link to the parent pointer and its two children.

#### 3.1 Node Class

---

**Algorithm 1** Node Class

---

- 1: Entry *leftEntry, rightEntry*
  - 2: Node *leftChild, rightChild, parent*
- 

#### Explanation

- *parent* - a link to parent node of type “Node”. It will be null, if the parent is not present i.e. the current node is the root itself.
- *leftChild* - stores the link to the left child is of type “Node”. It will be null if leftChild is not present. This can happen if current node is leaf node or tree has some skewness which will be handled by `compression()`
- *rightChild* - stores the link to the right child of type “node”. It will be null if rightChild is not present. This can happen if current node is leaf node or tree has some skewness which will be handled by `compression()`
- *leftEntry* - In the case of an internal node, *leftEntry* represents the minimum bounding rectangle containing all the nodes present in the left subtree of the current node. In the case of a leaf node, it will be a spatial point.
- *rightEntry* - In the case of an internal node, *rightEntry* represents the minimum bounding rectangle containing all the nodes present in the right subtree of the current node. In the case of a leaf node, it will be a spatial point.

#### 3.2 Entry Class

---

**Algorithm 2** Entry Class

---

- 1: Point *lowerBottom, upperTop*
  - 2: boolean *mark*
- 

#### Explanation

- *lowerBottom* - In case of leaf node, represents the spatial point. In case of an internal node, it represents the lower bottom coordinates of the minimum bounding rectangle.
- *upperTop* - In case of leaf node, it is null (to detect the difference between internal and leaf node). In case of an internal node, it represents the upper top coordinates of the minimum bounding rectangle.
- *mark* - By default, this value is false. *true* value denotes a new point has been physically added but it might not be logically added. Hence, the new point will be ignoolive by other threads/operations. Once, the new point becomes reachable from *root*, *mark* will be set to *false*.

#### 3.3 Atomic Referencing

For Lock-free implementation, we created root with the atomic reference class as it provides a flexible way to update values without use of synchronization. Moreover, we used `AtomicReferenceFeildUpdator` on the *leftChild* and *rightChild* of Node class to update the child links between parent and its children as it provides `AtomicReferenceFieldUpdater` provides us the capability of performing an atomic Compare-and-swap on a particular field of an object.

We have created two different instances of `atomicReferenceFieldUpdater`, one for each of the child node links.

Also, root is the only class based reference that needs to support atomic operations. Hence, it is created using the `AtomicReference` class which provides the capability to perform an atomic Compare-and-swap operation on the root node.

---

**Algorithm 3** AtomicReferencing

---

```
1: AtomicReference<Node> root = new AtomicReference<Node>()
2: AtomicReferenceFieldUpdater<Node,Node> leftChildUpdater = newUpdater(Node.class, Node.class, "left-Child")
3: AtomicReferenceFieldUpdater<Node,Node> rightChildUpdater = newUpdater(Node.class, Node.class, "rightChild")
```

---

## 4 Pseudo Code for Lock free implementation of R-tree

---

**Algorithm 4** Add ( *newPoint* )

---

```
1: restartAddition  $\leftarrow$  true
2: while restartAddition do
3:   if newPoint in tree then
4:     return
5:   end if
6:   if root is null then
7:     Create newNode with newPoint as an entry and its mark set as true
8:     if root.CAS(null, newNode) then
9:       restartAddition  $\leftarrow$  false
10:    else
11:      continue
12:    end if
13:  end if
14:  currNode  $\leftarrow$  root
15:  parent  $\leftarrow$  null
16:  parentChildLinkLeft  $\leftarrow$  true  $\triangleright$  Store the link of currNode to its parent i.e. whether it is left child or right child to its parent
17:  traversal  $\leftarrow$  true
18:  while traversal do
19:    if currNode is empty leaf then
20:      Make newNode a copy of currNode
21:      Add the newPoint in the empty slot of newNode with mark as true
22:    end if
23:    if currNode is an internal node then
24:      if checkAndCompressEmptyInternalNodes(currNode, parent, parentChildLinkLeft) then
25:        traversal  $\leftarrow$  false
26:        break
27:      end if
28:      Find minimum MBR expansion while adding the newPoint to left and right child
29:      if minimum MBR expansion is towards left then
30:        Traverse left
31:      else
32:        Traverse right
33:      end if
34:    end if
35:    if currNode is full leaf then
36:      Split the currNode into two entries
37:      Combine three entries i.e. two entries from currNode & one from newPoint into two diff nodes.
38:      Preserve the mark of old entries and set the mark of new entry as true
39:      Make newNode parent to these two different nodes.
40:    end if
41:    if currNode is full leaf or empty leaf then
42:      if currNode is root then
43:        if root.CAS(currNode, newNode) then
44:          restartAddition  $\leftarrow$  false
```

---

---

```

45:         else
46:             continue
47:         end if
48:     end if
49:     if currNode is its parent's leftChild then
50:         if leftChildUpdater.CAS(parent, currNode, newNode) then
51:             restartAddition  $\leftarrow$  false
52:         else
53:             if rightChildUpdater.CAS(parent, currNode, newNode) then
54:                 restartAddition  $\leftarrow$  false
55:             end if
56:         end if
57:     end if
58:     traversal  $\leftarrow$  false
59: end if
60: if restartAddition is false then
61:     updateMBR(currNode)
62:     unmark(newPoint)
63: end if
64: end while
65: end while

```

---



---

**Algorithm 5** *Unmark ( AddedPoint )*

---

```

1: restartUnmark  $\leftarrow$  true
2: while restartUnmark do
3:     currNode  $\leftarrow$  root
4:     newNode  $\leftarrow$  null
5:     didCompression  $\leftarrow$  false
6:     parentChildLink  $\leftarrow$  false
7:     foundPoint  $\leftarrow$  false
8:     parent  $\leftarrow$  null
9:     isParentChildLinkLeft  $\leftarrow$  true
10:    Make a queue Queue and add root to it, for traversal
11:    Also make a queue parentLinks and store link values while traversal
12:    while Queue is not empty do
13:        currNode  $\leftarrow$  Queue.poll()
14:        if currNode is null then
15:            break;
16:        end if
17:        if currNode is empty Leaf or full Leaf then
18:            if addedpoint is present in currNode then
19:                foundPoint  $\leftarrow$  true
20:            end if
21:        else
22:            if left child of currNode is not null then
23:                Add left child of currNode to Queue
24:            end if
25:            if right child of currNode is not null then
26:                Add right child of currNode to Queue
27:            end if
28:        end if
29:    end while
30:    if foundPoint is false then
31:        break
32:    end if
33:    if currNode is empty Leaf or full Leaf then
34:        Make the newNode a copy of currNode

```

---

▷ This means *currNode* is an internal Node

---

```

35:     if curr.leftEntry is addedPoint then
36:         newNode.leftEntry.mark  $\leftarrow$  true
37:     else curr.rightEntry is addedPoint
38:         newNode.rightEntry.mark  $\leftarrow$  true
39:     end if
40: end if
41: if currNode is the root then
42:     if root.CAS(currNode, newNode) then
43:         restartUnmark  $\leftarrow$  false
44:         break
45:     else
46:         continue
47:     end if
48: end if
49: Check the link of currNode to its parent i.e. whether it is left child or right child
50: if currNode is leftChild of its parent then
51:     if leftChildUpdater.CAS(parent, currNode, newNode) then
52:         restartUnmark  $\leftarrow$  false
53:         break
54:     else
55:         continue
56:     end if
57: end if
58: if currNode is rightChild of its parent then
59:     if rightChildUpdater.CAS(parent, currNode, newNode) then
60:         restartUnmark  $\leftarrow$  false
61:         break
62:     else
63:         continue
64:     end if
65: end if
66: end while

```

---



---

#### Algorithm 6 Delete ( *delPoint* )

---

```

1: restartDeletion  $\leftarrow$  true
2: while restartDeletion do
3:     if delPoint not in tree then
4:         return
5:     end if
6:     currNode  $\leftarrow$  root
7:     newNode  $\leftarrow$  null
8:     didCompression  $\leftarrow$  false
9:     parentChildLink  $\leftarrow$  false
10:    foundPoint  $\leftarrow$  false
11:    parent  $\leftarrow$  null
12:    Make a queue Queue and add root to it, for traversal
13:    while Queue is not empty do
14:        currNode  $\leftarrow$  Queue.poll()
15:        if currNode is null then
16:            break;
17:        end if
18:        if currNode is empty Leaf or full Leaf then
19:            if delpoint is present in currNode and its corresponding mark is false then
20:                foundPoint  $\leftarrow$  true
21:            end if
22:        else ▷ This means currNode is an internal Node

```

---

---

```

23:         if left child of currNode is not null then
24:             Add left child of currNode to Queue
25:         end if
26:         if right child of currNode is not null then
27:             Add right child of currNode to Queue
28:         end if
29:     end if
30: end while
31: if foundPoint is false then
32:     break
33: end if
34: Check if the currNode is empty leaf or full leaf
35: Make a newNode which will be used to replace the currNode using CAS
36: if currNode is a full leaf then
37:     Make the newNode a copy of currNode
38:     Make the entry that contains the delPoint as null
39: end if
40: if currNode is a empty leaf then Assign the newNode as null
41: end if
42: if currNode is the root then
43:     if root.CAS(currNode, newNode) then
44:         restartDeletion  $\leftarrow$  false
45:         break
46:     else
47:         continue
48:     end if
49: end if
50: Check the link of currNode to its parent i.e. whether it is left child or right child
51: if currNode is leftChild of its parent then
52:     if leftChildUpdater.CAS(parent, currNode, newNode) then
53:         restartDeletion  $\leftarrow$  false
54:         break
55:     else
56:         continue
57:     end if
58: end if
59: if currNode is rightChild of its parent then
60:     if rightChildUpdater.CAS(parent, currNode, newNode) then
61:         restartDeletion  $\leftarrow$  false
62:         break
63:     else
64:         continue
65:     end if
66: end if
67: end while

```

---



---

**Algorithm 7** Compression (*currNode*, *parent*, *parentChildLink*)

---

```

1: if currNode is null then
2:     return true
3: end if
4: leftChild  $\leftarrow$  currNode.leftChild
5: rightChild  $\leftarrow$  currNode.rightChild
6: if leftChild or rightChild are not null then
7:     return false
8: end if

```

---

---

```

9: if parent is null then
10:   root.CAS(currNode, null)
11: else
12:   if currNode is left child of parent then
13:     leftChildUpdater.CAS(parent, currNode, null)
14:   else  $\triangleright$  currNode is right child of parent
15:     rightChildUpdater.CAS(parent, currNode, null)
16:   end if
17: end if
18: return true

```

---



---

**Algorithm 8** RangeSearch(*Point* *p1*, *Point* *p2*)

---

```

1: Create a Entry with the given points, let it be range.
2: Create a hashSet of points, lets call it prevScan, which stores the points of a single scan
3: restartScan  $\leftarrow$  true
4: while restartScan do
5:   Create a hashSet of points, lets call it currScan, which stores the points of a single scan
6:   Make a queue Queue and add root to it, for traversal
7:   while Queue is not empty do
8:     currNode  $\leftarrow$  Queue.poll()
9:     if currNode is null then
10:       continue
11:     end if
12:     if currNode.leftEntry is not null then
13:       if currNode.leftEntry is a Point then
14:         if curr.leftEntry.lowerBottom is in range and curr.leftEntry.mark is false then
15:           Add curr.leftEntry.lowerBottom to currScan
16:         end if
17:       end if
18:     else
19:       if range and curr.leftEntry overlap then
20:         Add curr.leftChild to Queue
21:       end if
22:     end if
23:     if currNode.rightEntry is not null then
24:       if currNode.rightEntry is a Point then
25:         if curr.rightEntry.lowerBottom is in range and curr.rightEntry.mark is false then
26:           Add curr.rightEntry.lowerBottom to currScan
27:         end if
28:       end if
29:     else
30:       if range and curr.rightEntry overlap then
31:         Add curr.rightChild to Queue
32:       end if
33:     end if
34:   end while
35:   if prevScan is not null then
36:     if currScan and prevScan are same then
37:       We have found points for the current range, print these points.
38:       restartScan  $\leftarrow$  false
39:     end if
40:   end if
41:   prevScan  $\leftarrow$  currScan
42: end while

```

---

## 5 Correctness

- Addition

Once, we have physically added the node using CAS, we update the MBR for all its ancestors. After all the updates are completed, we begin to unmark the newly added node in the tree. Once the newly added point is unmarked, we consider our add operation to be completed. Therefore the linearization point of add is in function *unmark* line 42, 51 and 59. This is done to maintain linearizability of the entire algorithm. Moreover, the newly added node can only be deleted after it is unmarked i.e. the add operation has been successfully completed.

Also, the invariants are maintained after the addition is successfully completed. We split the leaf node if there is no space to add the new point. This adds a new level at that point. *updateMBR* makes sure the newly added point is reachable from the root node and the MBRs of internal nodes are up to date. We also check if an internal node has 0 children, we delete such an internal node through the *compression* algorithm.

- Deletion

For delete, the linearization point is at line 43, 52 and 60 of *delete* algorithm. This is where we use CAS to physically remove the point to be deleted.

We completely delete the leaf node when there is only one point in that leaf node which needs to be deleted. Thus, making sure there is no node with 0 entries.

- RangeSearch

For rangeSearch, we have linearization point at line 36 of *rangeSearch* algorithm. This is the point where the two sets are compared which shows the last two traversals have resulted into same result. To make sure the results of range search are linearizable, we have added the marking of newly added points. We only count a point in range search if it is not marked making sure that partially added nodes are not included.

## 6 Lock Freedom

We can observe that the CAS to add and delete is reattempted only if the node on which the operation is to be performed is modified. Before every reattempt of the CAS, the entire tree traversal happens which guarantees a fresh set of variables. Moreover, it also guarantees that a modify operation cannot take infinite number of steps without any modification to the data structure. It proves the lock-freedom for Add and Delete. In case of RangeSearch, we have used simple snapshot algorithm, it traverses the R tree twice, and if the traversal outputs the same result twice then we have the correct set of points for the range search. Since there are finite number of operations and threads, at some point the two traversal output are bound to match. Hence it guarantees lock-freedom in RangeSearch as well.

## 7 Amortized analysis

In Add and Delete, we have CAS operation during the traversal of the tree. In case of Add, the CAS failure can be caused when the parent node to which the new node is to be added is updated by some other thread. Therefore, it is necessary to traverse the tree again and find the new suitable addition point for the new node. In case of Delete, the CAS failure can be caused when the parent node of the node to be deleted is modified or the node which is to be deleted is modified by some other addition or deletion. Hence, in this scenario as well, it is necessary to traverse the tree again. If there are  $n$  nodes in the tree with a total of  $C_i$  number of concurrent operations, then the amortized number of steps per operation can be  $O(n * C_i)$ . Moreover, we can also say that this is asymptotically equivalent to  $O(n * C_p)$ , where  $C_p$  is the maximum number of concurrent operations at any point in the lifetime of our program.

## 8 Future Work

We can improve the rangeSearch algorithm by using better snapshot algorithms. Moreover, experimental analysis can also be done by comparing the lock-based and lock free implementations of the algorithm.

## Acknowledgments

This implementation of lock free R-tree is proposed as part of course project towards our course "Concurrent and Learned Data Structures", taught by Dr. Bapi Chatterjee at IIIT Delhi.



## References

- Chatterjee, Bapi. “Lock-Free Linearizable 1-Dimensional Range Queries.” Proceedings of the 18th International Conference on Distributed Computing and Networking, Association for Computing Machinery, 2017, pp. 1–10. ACM Digital Library, <https://doi.org/10.1145/3007748.3007771>.
- “R-Tree.” Wikipedia, 14 Mar. 2022. Wikipedia, <https://en.wikipedia.org/w/index.php?title=R-tree&oldid=1077176048>.
- Rtree-Chap1.Pdf. <https://tildesites.bowdoin.edu/~ltoma/teaching/cs340/spring08/Papers/Rtree-chap1.pdf>. Accessed 22 Apr. 2022.