
LOCK FREE MULTIDIMENSIONAL RANGE SEARCH

Aman Aggarwal
2018327
IIIT
Delhi
aman18327@iiitd.ac.in

Prasham Narayan
2018359
IIIT
Delhi
prasham18359@iiitd.ac.in

ABSTRACT

Multi-Dimensional Range Search is a fundamental problem in many application domains dealing with multidimensional data. R-tree is an ideal candidate for the multidimensional range search problem. However, there is no existing implementation of R-tree supporting concurrent operations. This paper introduces the lock-free concurrent R-tree, which implements an abstract data type (ADT) that provides the operations Add, Remove, RangeSearch. The operations in the Lock Free R-tree use single-word read and compare-and-swap (CAS) atomic primitives, which are readily supported on available multi-core processors. A lock-based implementation is also prepared as a benchmark.

Keywords Lock-Free · Lock-Based · R-tree · Range Search

1 Introduction

We propose using an external 2-way R-tree to implement the ADT = Add, Delete, RangeSearch in a 2-dimensional setting. An R-Tree is a tree-based data structure that can store spatial data indexes. Each entry in an internal node of an R-tree corresponds to MBR (Minimum bounding rectangle), which represents the smallest rectangle containing all its child nodes. R-trees use this concept of MBR to group nearby objects, which helps in spatial search. Each entry in the leaf nodes corresponds to a spatial point.

The main structure of our Lock free R tree is adapted from its sequential implementation. In the sequential implementation, we defined all the classes with its internal structure. Using the sequential implementation we also created a lock-based implementation of R tree. This lock-based implementation provided further insights on additional cases that are needed to be handled in concurrent setting.

Moreover, we will also compare the performance of lock-based implementation of R tree with its lock free implementation. Github Link - github.com/aman-agg/CLDS-Project

2 Abstract Data Type

Multi Dimensional Range Search Tree (MDRST) = {Add, Delete, RangeSearch}

- **Add** - adds point p to the tree
- **Delete** - deletes point p from the tree
- **RangeSearch** - Given a 2-dimensional range, return points lying inside that range.

2.1 Invariants

- Two way R tree i.e. it has two children (leftChild and rightChild).
- There will be atleast a single entry in a node.
- All internal nodes will have both the entries.
- MBR criteria i.e. parent node's MBR will envelop the MBR of its children.
- Tree is imbalanced.

3 Structure of Components

We will be using a 2-way R-tree. In a 2-way R-tree, each node consists of 2 entries (left entry and right entry). Each node contains a link to the parent pointer and its two children.

3.1 Node Class

Algorithm 1 Node Class

- 1: Entry *leftEntry, rightEntry*
 - 2: Node *leftChild, rightChild, parent*
-

Explanation

- parent - a link to parent node of type “Node”. It will be null, if the parent is not present i.e. the current node is the root itself.
- leftChild - stores the link to the left child is of type “Node”. It will be null if leftChild is not present. This can happen if current node is leaf node or tree has some skewness which will be handled by compression()
- rightChild - stores the link to the right child of type “node”. It will be null if rightChild is not present. This can happen if current node is leaf node or tree has some skewness which will be handled by compression()
- leftEntry - In the case of an internal node, leftEntry represents the minimum bounding rectangle containing all the nodes present in the left subtree of the current node. In the case of a leaf node, it will be a spatial point.
- rightEntry - In the case of an internal node, rightEntry represents the minimum bounding rectangle containing all the nodes present in the right subtree of the current node. In the case of a leaf node, it will be a spatial point.

3.2 Entry Class

Algorithm 2 Entry Class

- 1: Point *lowerBottom, upperTop*
-

Explanation

- lowerBottom - In case of leaf node, represents the spatial point. In case of an internal node, it represents the lower bottom coordinates of the minimum bounding rectangle.
- upperTop - In case of leaf node, it is null (to detect the difference between internal and leaf node). In case of an internal node, it represents the upper top coordinates of the minimum bounding rectangle.

3.3 Atomic Referencing

For Lock-free implementation, we created root with the atomic reference class as it provides a flexible way to update values without use of synchronization. Moreover, we used AtomicReferenceFeildUpdater on the leftChild and rightChild of Node class to update the child links between parent and its children as it provides AtomicReferenceFieldUpdater provides us the capability of performing an atomic Compare-and-swap on a particular field of an object.

We have created two different instances of atomicReferenceFieldUpdater, one for each of the child node links.

Also, root is the only class based reference that needs to support atomic operations. Hence, it is created using the AtomicReference class which provides the capability to perform an atomic Compare-and-swap operation on the root node.

Algorithm 3 AtomicReferencing

- 1: AtomicReference<Node> *root* = new AtomicReference<Node>()
 - 2: AtomicReferenceFieldUpdater<Node,Node> *leftChildUpdater* = newUpdater(Node.class, Node.class, "leftChild")
 - 3: AtomicReferenceFieldUpdater<Node,Node> *rightChildUpdater* = newUpdater(Node.class, Node.class, "rightChild")
-

4 Pseudo Code for Lock free implementation of R-tree

Algorithm 4 Add (*newPoint*)

```
1: restartAddition  $\leftarrow$  true
2: while restartAddition do
3:   if newPoint in tree then
4:     return
5:   end if
6:   if root is null then
7:     Create newNode with newPoint as an entry
8:     if root.CAS(null, newNode) then
9:       restartAddition  $\leftarrow$  false
10:    else
11:      continue
12:    end if
13:  end if
14:  currNode  $\leftarrow$  root
15:  parent  $\leftarrow$  null
16:  Store the link of currNode to its parent i.e. whether it is left child or right child to its parent
17:  traversal  $\leftarrow$  true
18:  while traversal do
19:    if currNode is empty leaf then
20:      Make newNode a copy of currNode
21:      Add the newPoint in the empty slot of newNode
22:    end if
23:    if currNode is an internal node then
24:      Find minimum MBR expansion while adding the newPoint to left and right child
25:      if minimum MBR expansion is towards left then
26:        Traverse left
27:      else
28:        Traverse right
29:      end if
30:    end if
31:    if currNode is full leaf then
32:      Split the currNode into two entries
33:      Combine three entries i.e. two entries from currNode & one from newPoint into two diff nodes.
34:      Make newNode parent to these two different nodes.
35:    end if
36:    if currNode is full leaf or empty leaf then
37:      if currNode is root then
38:        if root.CAS(currNode, newNode) then
39:          restartAddition  $\leftarrow$  false
40:        else
41:          continue
42:        end if
43:      end if
44:      if currNode is its parent's leftChild then
45:        if leftChildUpdater.CAS(parent, currNode, newNode) then
46:          restartAddition  $\leftarrow$  false
47:        else
48:          if rightChildUpdater.CAS(parent, currNode, newNode) then
49:            restartAddition  $\leftarrow$  false
50:          end if
51:        end if
52:      end if
53:      traversal  $\leftarrow$  false
54:    end if
```

```

55:      if restartAddition is false then
56:          updateMBR(currNode)
57:      end if
58:  end while
59: end while

```

Algorithm 5 Delete (*delPoint*)

```

1: restartDeletion  $\leftarrow$  true
2: while restartDeletion do
3:     if delPoint not in tree then
4:         return
5:     end if
6:     currNode  $\leftarrow$  root
7:     newNode  $\leftarrow$  null
8:     didCompression  $\leftarrow$  false
9:     parentChildLink  $\leftarrow$  false
10:    foundPoint  $\leftarrow$  false
11:    parent  $\leftarrow$  null
12:    Make a queue Queue and add root to it, for traversal
13:    while Queue is not empty do
14:        currNode  $\leftarrow$  Queue.poll()
15:        if currNode is null then
16:            break;
17:        end if
18:        if currNode is Internal node then
19:            Update parentChildLink
20:            Update parent
21:            if checkAndCompressSkewed(currNode, parent, parentChildLink) then
22:                didCompression  $\leftarrow$  true
23:                break
24:            else
25:                if left child of currNode is not null then
26:                    Add left child of currNode to Queue
27:                end if
28:                if right child of currNode is not null then
29:                    Add right child of currNode to Queue
30:                end if
31:            end if
32:        else
33:            if delPoint is present in currNode then
34:                foundPoint  $\leftarrow$  true
35:                break
36:            end if
37:        end if
38:    end while
39:    if didCompression then
40:        continue
41:    end if
42:    if foundPoint is false then
43:        break
44:    end if
45:    Check if the currNode is empty leaf or full leaf
46:    Make a newNode which will be used to replace the currNode using CAS
47:    if currNode is a full leaf then
48:        Make the newNode a copy of currNode
49:        Make the entry that contains the delPoint as null
50:    end if

```

```

51:  if currNode is a empty leaf then Assign the newNode as null
52:  end if
53:  if currNode is the root then
54:    if root.CAS(currNode, newNode) then
55:      restartDeletion  $\leftarrow$  false
56:      break
57:    else
58:      continue
59:    end if
60:  end if
61:  Check the link of currNode to its parent i.e. whether it is left child or right child
62:  if currNode is leftChild of its parent then
63:    if leftChildUpdater.CAS(parent, currNode, newNode) then
64:      restartDeletion  $\leftarrow$  false
65:      break
66:    else
67:      continue
68:    end if
69:  end if
70:  if currNode is rightChild of its parent then
71:    if rightChildUpdater.CAS(parent, currNode, newNode) then
72:      restartDeletion  $\leftarrow$  false
73:      break
74:    else
75:      continue
76:    end if
77:  end if
78: end while

```

Algorithm 6 Compression (*currNode*, *parent*, *parentChildLink*)

```

1:  if currNode is null then
2:    return true
3:  end if
4:  leftChild  $\leftarrow$  currNode.leftChild
5:  rightChild  $\leftarrow$  currNode.rightChild
6:  if leftChild and rightChild are not null then
7:    return false
8:  end if
9:  if leftChild or rightChild is null then
10:   if parent is null then
11:     root.CAS(currNode, null)
12:   else
13:     if currNode is left child of parent then
14:       leftChildUpdater.CAS(parent, currNode, null)
15:     end if
16:     if currNode is right child of parent then
17:       rightChildUpdater.CAS(parent, currNode, null)
18:     end if
19:   end if
20:   return true
21: end if

```

Algorithm 7 RangeSearch

To be added

5 Future Work

We are currently at the stage of integrating the compression algorithm into the add and delete algorithm of our R-tree. We are also working on implementing the lock free rangeSearch operation using wait free snapshot technique.

Acknowledgments

This implementation of lock free R-tree is proposed as part of course project towards our course "Concurrent and Learned Data Structures", taught by Dr. Bapi Chatterjee at IIIT Delhi. We had multiple discussions with him where he guided us and provided us multiple resources to tackle our problems. This project is only possible because of his continuous contributions.

References

- Chatterjee, Bapi. "Lock-Free Linearizable 1-Dimensional Range Queries." Proceedings of the 18th International Conference on Distributed Computing and Networking, Association for Computing Machinery, 2017, pp. 1–10. ACM Digital Library, <https://doi.org/10.1145/3007748.3007771>.
- "R-Tree." Wikipedia, 14 Mar. 2022. Wikipedia, <https://en.wikipedia.org/w/index.php?title=R-tree&oldid=1077176048>.
- Rtree-Chap1.Pdf. <https://tildesites.bowdoin.edu/~ltoma/teaching/cs340/spring08/Papers/Rtree-chap1.pdf>. Accessed 22 Apr. 2022.